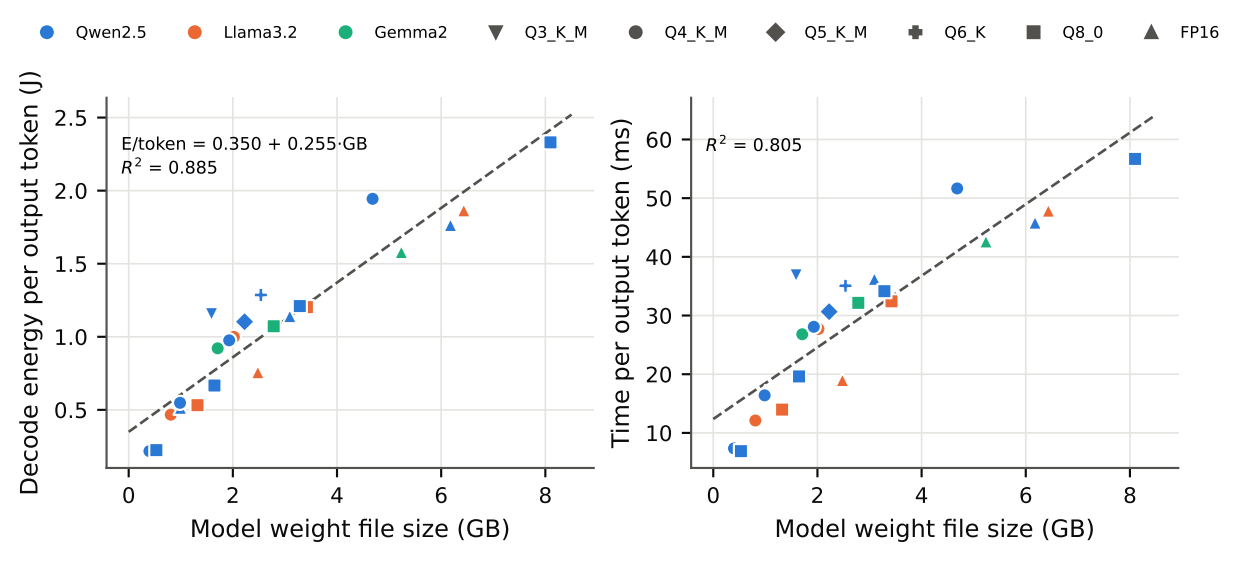
Qwen2.5
Llama3.2
Gemma2
Q3_K_M
Q4_K_M
Q5_K_M
Q6_K
Q8_0
FP16
E/token = 0.350 + 0.255·GB
R^2 = 0.885
Decode energy per output token (J)
Model weight file size (GB)
R^2 = 0.805
Time per output token (ms)
Model weight file size (GB)